

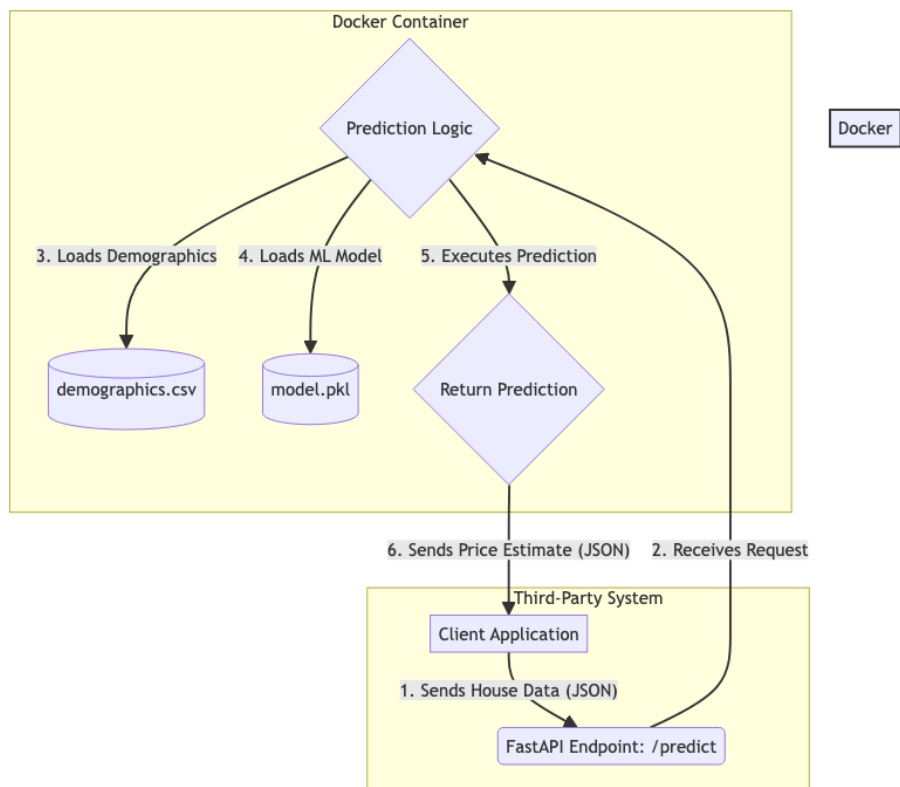
Technical Presentation: House Price Prediction API

1. Project Overview

Deployed a machine learning model for house price prediction as a scalable RESTful service. Key deliverables: API, test script, and model performance evaluation.

2. Architecture and Design Choices

Simple, robust architecture using modern open-source tools.



- **API Framework: FastAPI**
 - High performance, asynchronous support, automatic validation, OpenAPI documentation.
- **Containerization: Docker**
 - Consistent environment for development/deployment, simplified dependency management.

- **Data Handling:**
 - House data sent via JSON to `/predict` endpoint.
 - `zipcode_demographics.csv` loaded into pandas DataFrame at startup.
 - **Production Consideration:** Migrate demographic data to PostgreSQL for scalability.
-

3. Model Evaluation and Improvement

Initial Model

- **Algorithm:** `KNeighborsRegressor` with `RobustScaler`.
- **Features:** Small subset of numeric columns.
- **Performance:**
 - **R-squared:** 0.7284
 - **Mean Absolute Error:** \$102,337.19

Improved Model

Developed an improved model per project recommendations.

- **Algorithm:** `GradientBoostingRegressor`
 - More powerful ensemble method with higher accuracy than k-NN.
- **Feature Engineering:**
 - Used all available numeric features.
 - Extracted `sale_year` and `sale_month` from `date` field.
- **Performance:**
 - **R-squared:** **0.8804** (+21% improvement)
 - **Mean Absolute Error:** **\$69,841.29** (32% improvement)

Substantial improvement in prediction accuracy.

4. Scalability and Future Work

- **Scalability:** Containerized API allows easy scaling. Kubernetes deploys multiple instances with load balancing for high availability.
- **Model Versioning:** Current setup loads model at startup. Robust strategy:
 1. Store model artifacts in MLflow or S3.
 2. Create `/-/reload_model` endpoint for zero-downtime updates.
- **Future Improvements:**
 - Hyperparameter tuning with `GridSearchCV`.
 - Advanced feature engineering for geographical data.

5. Cloud Deployment Strategy

Production deployment in the cloud using:

Container Orchestration

- **Kubernetes** (preferred) or **AWS ECS**
 - Horizontal scaling based on demand
 - High availability through replica management
 - Rolling updates for zero-downtime deployments
 - Key Components: Deployments, Services, Ingress, ConfigMaps, Secrets

Cloud Provider Options

- **AWS:** EKS, ALB, S3, CloudWatch
- **GCP:** GKE, Cloud Load Balancing, Cloud Storage, Cloud Monitoring
- **Azure:** AKS, Application Gateway, Azure Blob Storage, Azure Monitor

Model Registry and Versioning

- **MLflow** (recommended):
 - Model registry with version control
 - Performance metrics comparison
 - Model stage transitions
- **Alternatives:** S3/GCS/Azure Blob with manual versioning, DVC, Weights & Biases

CI/CD Pipeline

- **GitHub Actions** or **GitLab CI/CD**:
 - Stages: Testing, Building, Security Scanning, Deployment
 - Tools: pytest, Docker Hub/ECR, kubectl, notifications

Monitoring and Observability

- **APM:** Prometheus, Grafana
- **Metrics:** Response times, throughput, error rates, resource utilization
- **Logging:** ELK Stack or similar solutions
- **Health Checks:** Liveness/readiness probes, downtime alerts

API Authentication and Security

- **Methods:** JWT, API Keys, OAuth2
- **Implementation:** FastAPI dependency injection, rate limiting
- **Best Practices:** TLS encryption, security scanning, least privilege

Auto-scaling

- Kubernetes Horizontal Pod Autoscaler (HPA)
- Scale based on CPU utilization or custom metrics (requests/second, response time)

Additional Considerations

- **Backup/Recovery:** Regular backups, cross-region replication
- **Cost Optimization:** Resource tuning, spot instances, auto-scaling policies

Comprehensive cloud deployment strategy ensures scalability, reliability, and maintainability.